

EN.600.461 Computer Vision

Final Project

Recognizing and Translating Text from Images

Joon Hyuck (James) Choi
Senior Undergraduate
The Johns Hopkins University
3400 N Charles Street
Baltimore, MD 21218, USA
jchoi100@jhu.edu

Abstract

This work explores optical character recognition (OCR) in photos of typed and hand-written documents. It first explores basic preprocessing of photos of documents using `OpenCV`¹ functionalities for blurring, thresholding, and denoising. It next discusses the use of `tesseract-ocr`² to perform OCR. We then discuss the incorporation of the Google Cloud Translation API³ to translate the OCR results to different languages. It finally discusses template matching using the Johns Hopkins University (JHU) logo and signatures of three U.S. presidents on sample document images.

1 Introduction

This work achieved the following main goals stated in the original proposal: 1) Given a photo of a document, convert it into a clean scanned version; 2) Take the scanned version and perform OCR. Optionally, this work took the OCR output and translated the text into different languages using the Google Cloud Translation API. This work also experimented with template matching using the JHU logo in order to determine whether the input document is an official JHU document or not. Lastly, this work performed CNN training on the MNIST handwritten digit dataset as an experiment.

¹<http://opencv.org/>

²<https://github.com/tesseract-ocr>

³<https://cloud.google.com/translate/>

2 Methods

In this section, we discuss the methods we took and external libraries used for each stage of our work.

2.1 Image Preprocessing

We implemented our program in Python using `OpenCV`. In order to feed the OCR algorithm clean input to achieve best performance, we used `cv2.medianBlur` to smoothen the input image with an aperture size of 5. Then, we performed binary thresholding on the smoothened image using `cv2.adaptiveThreshold` with *adaptiveMethod*= `ADAPTIVE_THRESH_GAUSSIAN_C`, *threshold* = *binary*, *blockSize*=`5x5`, and *C*=2. The adaptive method we selected uses the weighted sum of the *blockSize* x *blockSize* neighborhood of pixel (*x*, *y*) - *C* as its threshold value. Parameter values were chosen empirically.

Finally, we performed denoising on the thresholded image to remove noise and make the output clean. We used `cv2.fastNlMeansDenoising` with *templateWindowSize*=7, *searchWindowSize* = 21, and *h*=7. Parameters here were also chosen empirically. The motivation behind denoising in our work was that document scans usually have many random specs spread throughout. Such noise can greatly reduce OCR accuracy because an OCR algorithm may confuse a random spec with punctuation marks or associate a spec with an actual character near by (e.g. confuse an *l* and ‘ with an *i*).

2.2 OCR

In order to perform OCR on the preprocessed inputs, we used the `tesseract-ocr` library

through its python wrapper `pytesseract`. We used the raw output from `pytesseract`'s `image_to_string` method to pass into python's file writer and Google Cloud Translation API.

As a side experiment, we used the `Keras`⁴ library to train the MNIST⁵ dataset of 70,000 hand-written digits on three different convolutional neural networks (CNN). We referred to code available online⁶ to build the three CNNs.

2.3 Text Translation

We made use of Google Cloud's Translation API in order to translate the OCR processed document into several different languages based on user command line input. We referred to external code⁷ and modified the details to fit our purpose.

2.4 Template Matching

Aside from OCR, which was the main focus of this project, we experimented with template matching in images. We tested with various targets. First, we targeted to identify signatures of U.S. Presidents Barack Obama, George W. Bush, and Bill Clinton in document images. Our dataset consisted of hand-written letters written by the three presidents as input images and their respective signature images as templates. Performance for this experiment was poor, and we decided to experiment with an easier, fixed template. Potential reasons for low performance are described below.

After experimenting with signatures of U.S. presidents, we used the JHU logo to determine whether the input document was an official JHU document or not. We simply assumed that official JHU documents contained a JHU logo for our experimental purposes. We used `cv2.matchTemplate` with our input document as the source image, multiple JHU logo images (each with different sizes since `cv2.matchTemplate` is scale sensitive) as templates, and matching method `cv2.TM_CCOEFF_NORMED`. If the normalized `TM_CCOEFF` value came out to be greater than 0.51

for any of the input JHU logo templates of various sizes, we determined that the input document contained a JHU logo in it. The matching method and determinant value of 0.51 were chosen empirically.

Performance for the JHU logo template experiments was better compared to that of the U.S. president signature experiments. The reason that the task is harder for the earlier experiment is that the JHU logo is fixed in terms of how it is comprised (lines, curves, color density, etc.) whereas signatures vary each time it is signed (penstroke width, curvature, ratio of one part of the signature to another, etc.). Therefore, we think that the "signature matching task" bears more resemblance to an object recognition task than an OCR task.

3 Results

Here, we qualitatively discuss OCR results on typed and handwritten input documents.

3.1 Typed Documents

We ran our code on seven document images with different font types and font sizes. The images were `cv_proj_description.png`, `cv_jchoi_proposal.png`, `aesop_fox.png`, `lav_notes1.png`, `aesop_miller.png`, `aesop_goose.png`, and `aesop_cat.png`.

With the exception of `aesop_goose.png`, performance on all documents were nearly perfect. A sample input image is shown in Figure 1. The reason that `aesop_goose.png` did not produce satisfactory results was that the font in the document was relatively thick and bold. Therefore, after the preprocessing stage, the alphabet characters in the thresholded document were hollow with just the edges remaining. Therefore, the OCR algorithm could not recognize any of the characters in this text.

The two images `cv_jchoi_proposal.png` and `lav_notes1.png` were taken under sub-optimal conditions: irregular lighting, partial skew, and/or glossy surfaces. We ran our code on these two document images. Moreover, these documents had mixed font types and sizes, and the structure of the document was more complicated than the previous inputs. Despite the less optimal settings, we correctly denoised and thresholded the images and got similar results as before.

⁴<https://keras.io/>

⁵<http://yann.lecun.com/exdb/mnist/>

⁶<http://machinelearningmastery.com/handwritten-digit-recognition-using-convolutional-neural-networks-python-keras/>

⁷<http://github.com/mouuff/mtranslate>

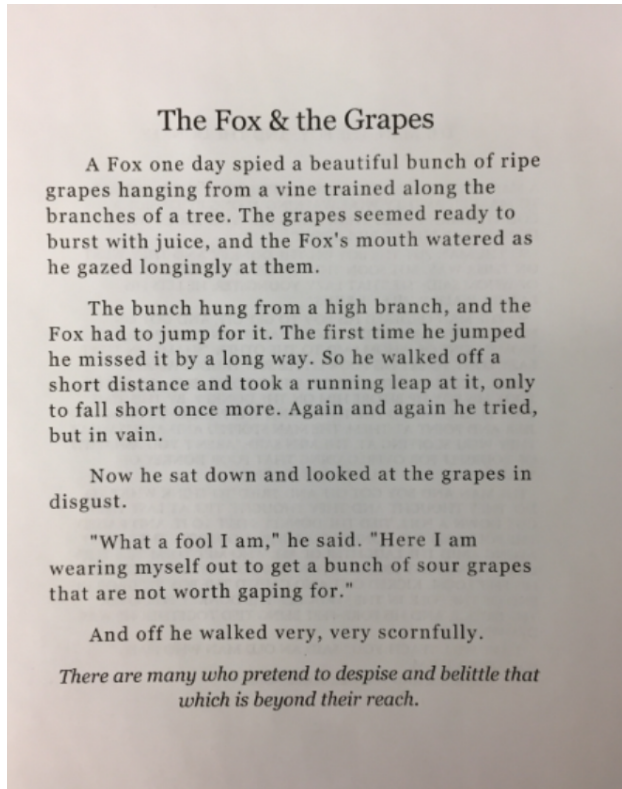


Figure 1: Sample typed document aesop_fox.png

3.2 Handwritten Documents

We applied our code on photos of handwritten documents. Performance on very neatly written documents (hand_fox.png and hand_goose.png) was nowhere close to what we saw for the printed dataset. Moreover, any reasonably neatly handwritten documents (hand_miller.png and hand_cat.png) showed even poorer performance. Figure 2 shows one of our reasonably neat images hand_cat.png. Furthermore, our code was not able to detect a single word from document images with much less neat handwriting. Such images were obama1.png, obama2.png, obama3.png, bush1.png, bush2.png, clinton1.png, and clinton2.png. Figure 3 shows obama1.png.

4 Conclusion

We explored optical character recognition on photographed documents in various settings and formats. We tested with optimal and suboptimal photography settings, experimented with documents us-

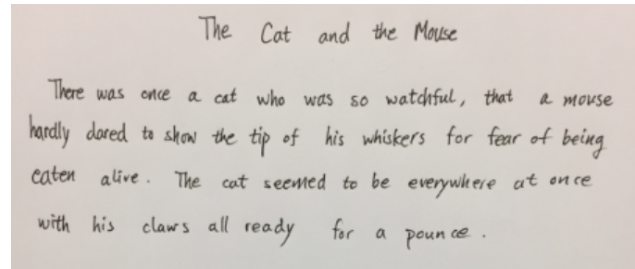


Figure 2: Reasonable handwritten doc hand_cat.png

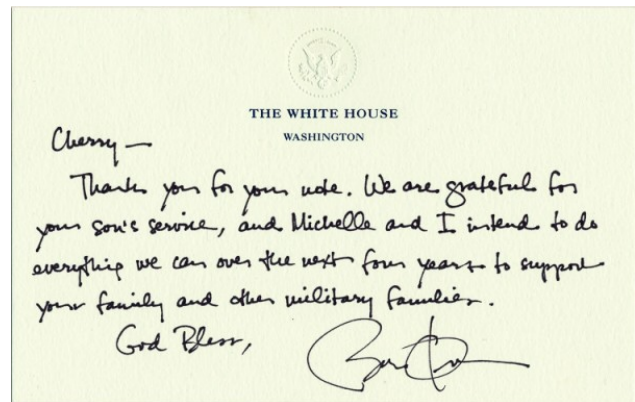


Figure 3: Less neatly handwritten doc obama1.png

ing various fonts, and compared performance on printed and handwritten documents. For printed documents, we saw highly accurate results under both optimal and suboptimal settings, even if they contained various font sizes and font types. However, we were not able to achieve accurate results with handwritten document photos, disregarding the neatness of the handwriting. For sufficiently small portions of handwritten documents (e.g. a single word or a short sentence), the OCR algorithm was able to produce correct output occasionally. However, even this was not often enough to be considered reliable.

Nonetheless, on typed documents, we were able to effectively preprocess the document image and perform OCR on the cleaned data. Moreover, we leveraged the Google Cloud Translation API to translate documents into several different languages. Our final experiments on template matching also deserves future development into a project of its own.

How to Run the Code

Please extract `FINAL_PROJECT-jchoi100.zip` to view and run the code.

Assuming that the user has installed Python 2.5+, OpenCV, pytesseract, and tesseract (with dependencies), run the following command in the same directory as `driver.py`.

Command:

```
$python driver.py {path-to-image-  
file} {list of target languages  
separated by single space}
```

Sample usage: \$python driver.py
data/doc7.png es fr ko

The sample command above performs OCR on `doc7.png`, outputs the thresholded image, the text in the original input language, and the text translated in Spanish, French, and Korean.

Saved Output from Run Experiments

The output that we created by running our code on all the images in our dataset are saved under `FINAL_PROJECT-jchoi100/code/data`.

Acknowledgments

We thank the image processing and OCR libraries (OpenCV and tesseract), the code we referred to for constructing CNNs using Keras, MNIST for the handwritten digit dataset, and Google for the translation API.

And thanks to Professor Reiter and the TAs for a great semester! Have a great winter break.